

cnn-and-ann

June 4, 2024

1 Convolution Neural Network (CNN) & Artificial Neural Network (ANN)

Name: Andrea Filiberto Lucas ID No: 0279704L Course: ARI2201 - Artificial Intelligence (AI)

1.1 Introduction

This notebook focuses on the in-depth analysis and classification of respiratory sound data. Respiratory sounds, such as wheezes and crackles, serve as significant indicators for various respiratory conditions. Analysing these sounds can help with the early detection and diagnosis of diseases such as COPD, asthma, and pneumonia. This approach uses machine learning (ML) techniques to improve diagnostic accuracy and efficiency in clinical settings.

1.1.1 Objectives

- **Data Loading & Preprocessing:** Load the respiratory sound files and corresponding diagnosis data, preprocess the data, and extract relevant features.
- **Data Visualization:** Visualize class distributions, audio file durations, and other relevant metrics to understand the dataset better.
- **Model Building and Evaluation:** Build and train a Convolutional Neural Network (CNN) and an Artificial Neural Network (ANN) to classify the respiratory sounds into different disease categories. Evaluate the model using various metrics to ensure its performance.
- **Result Interpretation:** Analyze the results, visualize model performance, and interpret the findings to gain insights into the model's effectiveness.

1.2 Convolutional Neural Network (CNN)

A Convolutional Neural Network (CNN) is a specialized type of Artificial Neural Network (ANN) designed for processing structured grid data, such as images. CNNs are particularly effective for image recognition and classification tasks due to their ability to automatically and adaptively learn spatial hierarchies of features through convolutional layers, pooling layers, and fully connected layers.

1.3 Artificial Neural Network (ANN)

An Artificial Neural Network (ANN) is a broad category of machine learning models inspired by the human brain's neural networks. ANNs consist of interconnected layers of nodes (neurons), including an input layer, one or more hidden layers, and an output layer. These networks can be trained to recognize patterns and make predictions based on input data, and they are used for a

wide range of tasks including regression, classification, and time series prediction. `##` Importing the necessary libraries & modules. First, all libraries used in this Jupyter notebook are imported:

- `os`, `os.path`, `time`: Directory operations, file status checks, and delays.
- `tabulate`, `tqdm`: Text tables and progress bars.
- `pandas`, `numpy`, `librosa`: Data manipulation, numerical operations, and audio processing.
- `matplotlib.pyplot`: Plotting and visualizations.
- `sklearn.model_selection`, `timeit`: Dataset splitting and time measurement.
- `wave`: Handling .wav audio files.
- `sklearn.metrics`, `seaborn`: Model evaluation and statistical visualization.
- `itertools.cycle`: Cycling through colors.
- `tensorflow.keras`: Building, regularizing, optimizing, and visualizing neural networks, including specific layers and the Adam optimizer.

```
[1]: # Importing the os module for directory operations
import os

# Importing isfile and join functions from os.path to check file status and
# create file paths
from os.path import isfile, join

# Import time for delays
import time

# Import tabulate for generating tables in text format
from tabulate import tabulate # type: ignore

# Import tqdm for creating progress bars
from tqdm import tqdm # type: ignore

# Importing pandas for data manipulation and analysis
import pandas as pd # type: ignore

# Importing librosa for audio processing
import librosa # type: ignore

# Importing pyplot from matplotlib for plotting
import matplotlib.pyplot as plt # type: ignore

# Importing numpy for numerical operations
import numpy as np # type: ignore

# Importing train_test_split for splitting datasets
from sklearn.model_selection import train_test_split # type: ignore

# Importing timer to measure time intervals
from timeit import default_timer as timer
```

```

# Importing the wave module to handle .wav audio files
import wave

'''
===== Evaluation =====
'''

# Importing classification_report, confusion_matrix, roc_curve, auc,
↳precision_recall_curve for model evaluation
from sklearn.metrics import classification_report, confusion_matrix, roc_curve,
↳auc, precision_recall_curve # type: ignore
# Importing seaborn for statistical data visualization
import seaborn as sn # type: ignore

# Import cycle from itertools for cycling through colors
from itertools import cycle

'''
===== TENSORFLOW =====
'''

# Importing Sequential and Model from tensorflow.keras for building neural
↳networks
from tensorflow.keras.models import Sequential, Model # type: ignore

# Importing regularizers and optimizers from tensorflow.keras for
↳regularization and optimization
from tensorflow.keras import regularizers, optimizers # type: ignore

# Importing various neural network layers from tensorflow.keras
from tensorflow.keras.layers import Dense, Conv1D, Flatten, Activation,
↳MaxPooling1D, Dropout # type: ignore

# Importing plot_model for model visualization and to_categorical for one-hot
↳encoding
from tensorflow.keras.utils import plot_model, to_categorical # type: ignore

# Import the Adam optimizer from tensorflow.keras
from tensorflow.keras.optimizers import Adam # type: ignore

```

```

/Users/afl/Library/Python/3.9/lib/python/site-packages/urllib3/__init__.py:35:
NotOpenSSLWarning: urllib3 v2 only supports OpenSSL 1.1.1+, currently the 'ssl'
module is compiled with 'LibreSSL 2.8.3'. See:
https://github.com/urllib3/urllib3/issues/3020
warnings.warn(

```

1.4 Parsing the Dataset

A `Diagnosis` class is defined with an initializer to store patient information:

- `patient_id`: The ID of the patient.
- `diagnosis`: The diagnosis of the patient.
- `audio_file_path`: The file path to the patient's audio file.

```
[2]: class Diagnosis:
      def __init__(self, patient_id, diagnosis, audio_file_path):
          self.patient_id = patient_id
          self.diagnosis = diagnosis
          self.audio_file_path = audio_file_path
```

1.4.1 Paths

Dynamic paths are defined for accessing files:

- `PATH`: Directory path to audio and text files.
- `diagnosis_file`: Path to the patient diagnosis CSV file.

```
[3]: # Dynamic paths
PATH = '../Respiratory_Sound_Database/audio_and_txt_files/'
diagnosis_file = '../Respiratory_Sound_Database/patient_diagnosis.csv'
```

1.4.2 Retrieving Audio Files from a Directory

`get_wav_files(PATH)`: - Checks if the directory exists, raises an error if not. - Retrieves and sorts `.wav` files from the directory. - Raises an error if no `.wav` files are found. - Returns the list of `.wav` files and the directory path.

```
[4]: # Function to get a list of .wav files in the audio directory
def get_wav_files(PATH):
    if not os.path.isdir(PATH):
        raise ValueError(f"The directory {PATH} does not exist.")

    try:
        # Get all files in the directory
        files = [f for f in os.listdir(PATH) if os.path.isfile(os.path.
↵join(PATH, f))]
        # Filter the files to only include .wav files
        wav_files = sorted([f for f in files if f.endswith('.wav')])
    except Exception as e:
        raise RuntimeError(f"Error while accessing files in the directory_
↵{PATH}: {e}")

    if not wav_files:
        raise FileNotFoundError("No .wav files found in the directory.")
```

```
return wav_files, PATH
```

1.4.3 Loading and Mapping Diagnosis Data to Audio Files

`load_diagnosis_data(PATH, diagnosis_file)`: - Checks if the diagnosis file exists, raises an error if not. - Reads the diagnosis file into a DataFrame, handling possible read errors. - Retrieves the list of .wav files and the audio path. - Creates a dictionary mapping patient IDs to diagnoses. - Iterates over .wav files, extracting patient IDs and associating diagnoses. - Creates `Diagnosis` objects for each .wav file, appending them to a list. - Returns the list of `Diagnosis` objects.

```
[5]: def load_diagnosis_data(PATH, diagnosis_file):
    # Check if diagnosis file exists
    if not os.path.isfile(diagnosis_file):
        raise ValueError(f"The file {diagnosis_file} does not exist.")

    try:
        # Read diagnosis file into a DataFrame
        diagnosis_df = pd.read_csv(diagnosis_file)
    except pd.errors.EmptyDataError:
        raise ValueError("Diagnosis file is empty.")
    except pd.errors.ParserError:
        raise ValueError("Diagnosis file is not properly formatted.")
    except Exception as e:
        raise RuntimeError(f"Error while reading the diagnosis file_{diagnosis_file}: {e}")

    if diagnosis_df.empty:
        raise ValueError("Diagnosis file is empty.")

    # Get the list of .wav files and the audio path
    wav_files, audio_path = get_wav_files(PATH)

    try:
        # Create a dictionary mapping patient IDs to diagnoses
        diagnosis_dict = {row.iloc[0]: row.iloc[1] for _, row in diagnosis_df.iterrows()}
    except KeyError as e:
        raise ValueError(f"Missing expected columns in diagnosis data: {e}")
    except Exception as e:
        raise RuntimeError(f"Error while processing the diagnosis data: {e}")

    diagnosis_list = []
    for idx, wav_file in enumerate(wav_files):
        try:
            # Extract the patient ID from the .wav file name
            patient_id = int(wav_file[:3])
```

```

        # Get the diagnosis for the patient ID from the dictionary, default
        ↪to "Unknown" if not found
        diagnosis = diagnosis_dict.get(patient_id, "Unknown")
        # Create a Diagnosis object and append it to the diagnosis_list
        diagnosis_list.append(Diagnosis(idx, diagnosis, os.path.
        ↪join(audio_path, wav_file)))
    except ValueError as e:
        raise ValueError(f"Error processing file {wav_file}: {e}")
    except Exception as e:
        raise RuntimeError(f"Unexpected error processing file {wav_file}:
        ↪{e}")

    return diagnosis_list

```

1.4.4 Comprehensive Analysis of Disease Distribution in Audio Files

The following code processes audio files and their corresponding diagnoses, generating various visualizations and statistics:

- **Constants Section:**
 - Lists audio files in the directory and extracts patient IDs.
 - Reads the patient diagnosis file and extracts labels for the audio files.
 - Computes and prints class counts.
- **Table of Classes:**
 - Converts class counts to a table and prints it.
- **Info Section:**
 - Calculates and prints the total number of classes and patients.
- **Bar Chart:**
 - Plots a bar chart showing the count of diseases in the sound files.
- **Pie Chart of Class Distribution:**
 - Plots a pie chart showing the distribution of classes in the sound files.
- **Histogram of Audio File Duration:**
 - Analyzes and plots a histogram of the durations of the audio files, including the mean duration.
- **Box Plot of Audio File Duration:**
 - Creates a box plot of the audio file durations, with annotations for outliers.
- **Statistical Summary of Audio File Durations:**
 - Computes and prints a statistical summary of the audio file durations.

```

[6]: try:
    ''' ===== CONSTANTS
    ↪===== '''

    # List audio files in the specified directory
    filenames = [f for f in os.listdir(PATH) if os.path.isfile(os.path.
    ↪join(PATH, f)) and f.endswith('.wav')]
    if not filenames:
        raise FileNotFoundError(f"No .wav files found in directory {PATH}")

```

```

# Extract patient IDs corresponding to each file
p_id_in_file = np.array([int(name.split('_')[0]) for name in filenames])

# Read patient diagnosis file
p_diag = pd.read_csv(diagnosis_file, header=None)
if p_diag.empty:
    raise ValueError(f"The diagnosis file {diagnosis_file} is empty or
↳improperly formatted.")

# Extract labels for audio files
labels = np.array([p_diag[p_diag[0] == x][1].values[0] for x in
↳p_id_in_file])
if labels.size == 0:
    raise ValueError("No matching patient IDs found in diagnosis file.")

# Get class counts
unique_elements, counts_elements = np.unique(labels, return_counts=True)
class_counts = np.asarray((unique_elements, counts_elements))

''' ===== TABLE of CLASSES
↳===== '''
# Convert class counts to a list of lists for tabulate
class_counts_list = class_counts.T.tolist()

# Print class counts in the form of a table
print(tabulate(class_counts_list, headers=['Classes', 'Count']))

time.sleep(1)

''' ===== INFO (print)
↳===== '''
# Calculate and print the total number of classes
total_classes = len(class_counts_list)
print('\n\033[1;92mTotal Number of Classes:', total_classes, '\033[0m')

time.sleep(1)

# Calculate and print the total number of patients
total_patients = np.sum(counts_elements)
print('\033[1;92mTotal Number of Patients:', total_patients, '\033[0m')

time.sleep(1)

''' ===== BAR CHART (Disease Count in Sound Files)
↳===== '''
# Define custom colors

```

```

colors = ['#ff9999', '#66b3ff', '#99ff99', '#ffcc99', '#c2c2f0']

# Plot class counts
plt.figure(figsize=(18, 12))
bars = plt.bar(unique_elements, counts_elements, align='center', alpha=0.7,
color=colors, edgecolor='black')

# Labels and title
plt.xlabel('Diseases', fontsize=14)
plt.ylabel('Number of Patients', fontsize=14)
plt.title('Disease Count in Sound Files (.wav)', fontsize=16)

# Rotate x-axis labels for better readability
plt.xticks(rotation=45, fontsize=12)
plt.yticks(fontsize=12)

# Calculate the upper limit for y-axis
y_max = max(counts_elements)
y_max = (y_max // 50 + 1) * 50 # Round up to the nearest multiple of 50

# Set custom y-axis range
plt.ylim(0, y_max)

# Set y-axis ticks to increment by 50
plt.yticks(np.arange(0, y_max + 1, 50))

# Add grid lines for better readability
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Add labels on top of each bar
for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width() / 2.0, height, f'{height}',
ha='center', va='bottom', fontsize=12)

# Show the plot
plt.show()

time.sleep(1)

''' ===== PIE CHART of CLASS DISTRIBUTION_
===== '''

# Generate pastel colors using seaborn
colors = sn.color_palette("pastel", len(unique_elements))

# Plot a pie chart of class distribution
plt.figure(figsize=(14, 8))

```



```

wedges, texts, autotexts = plt.pie(
    counts_elements, labels=unique_elements, autopct='%1.1f%%',
↪startangle=140,
    colors=colors, shadow=False, pctdistance=0.85, labeldistance=1.1
)

# Equal aspect ratio ensures that pie is drawn as a circle.
plt.axis('equal')

# Improve the appearance of text
for text in texts:
    text.set_fontsize(12)
for autotext in autotexts:
    autotext.set_fontsize(10)
    autotext.set_color('white')
    autotext.set_weight('bold')

# Add a legend
plt.legend(wedges, unique_elements, title="Classes", loc="center left",
↪bbox_to_anchor=(1, 0, 0.5, 1), fontsize=12)

# Add a title
plt.title('Class Distribution in Sound Files (.wav)', fontsize=16)

# Show the plot
plt.show()

time.sleep(1)

# Analyze the duration of each audio file
durations = []
for file in filenames:
    try:
        with wave.open(os.path.join(PATH, file), 'r') as wav_file:
            frames = wav_file.getnframes()
            rate = wav_file.getframerate()
            duration = frames / float(rate)
            durations.append(duration)
    except wave.Error as e:
        print(f"Error processing {file}: {e}")

# Plot a histogram of audio file durations
# Calculate mean duration
mean_duration = np.mean(durations)

''' ===== HISTOGRAM of AUDIO FILE DURATION_
↪===== '''

```

```

plt.figure(figsize=(14, 8))
counts, bins, patches = plt.hist(durations, bins=20, edgecolor='black',
↪alpha=0.7, color='skyblue', label='Durations')
plt.axvline(mean_duration, color='red', linestyle='dashed', linewidth=2,
↪label=f'Mean: {mean_duration:.2f} sec')

# Adding labels and title
plt.xlabel('Duration (seconds)', fontsize=14)
plt.ylabel('Number of Files', fontsize=14)
plt.title('Distribution of Audio File Durations', fontsize=16)

# Adding grid lines for better readability
plt.grid(True, linestyle='--', alpha=0.7)

# Customizing tick labels
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)

# Adding legend
plt.legend(fontsize=12)

# Adding annotation for the mean
plt.text(mean_duration + 0.5, max(counts) * 0.9, f'Mean: {mean_duration:.
↪2f} sec', fontsize=12)

# Adding count labels above each bar
for patch, count in zip(patches, counts):
    plt.text(patch.get_x() + patch.get_width() / 2, count, str(int(count)),
             ha='center', va='bottom', fontsize=10)

plt.show()

time.sleep(1)

''' ===== BOX PLOT of AUDIO FILE DURATION
↪===== '''
plt.figure(figsize=(14, 8))

# Creating a horizontal boxplot with enhanced features
box = plt.boxplot(durations, vert=False, patch_artist=True,
                  boxprops=dict(facecolor='lightblue', color='blue'),
                  whiskerprops=dict(color='blue'),
                  capprops=dict(color='blue'),
                  medianprops=dict(color='red'),
                  flierprops=dict(marker='o', color='red', alpha=0.5))

# Adding grid lines for better readability

```

```

plt.grid(True, linestyle='--', alpha=0.7)

# Adding labels and title
plt.xlabel('Duration (seconds)', fontsize=12)
plt.title('Box Plot of Audio File Durations', fontsize=14)
plt.xticks(fontsize=10)
plt.yticks([]) # Remove y-axis ticks as they are not needed for a single
↳boxplot

# Adding annotations for outliers
for flier in box['fliers']:
    for outlier in flier.get_ydata():
        plt.annotate(f'{outlier:.2f}', xy=(outlier, 1), xytext=(5, 0),
                      textcoords='offset points', fontsize=10, color='red')

plt.show()

time.sleep(1)

''' ===== STATISTICAL SUMMARY of AUDIO FILE
↳DURATIONS ===== '''
durations_series = pd.Series(durations)
summary_stats = durations_series.describe().round(2)

# Select the specific statistics to display
custom_summary = summary_stats[['mean', 'std', 'min', 'max']]

# Convert the summary to a DataFrame for tabulate
custom_summary_df = custom_summary.to_frame().transpose()

# Print the statistical summary using tabulate
print('\n\033[1;92mStatistical Summary of Audio File Durations (in
↳seconds)\033[0m')
print(tabulate(custom_summary_df, headers='keys', tablefmt='simple'))


except Exception as e:
    print("\033[1;91mWARNING:\033[0m An error occurred while processing class
↳counts:")
    print("Details:", e)

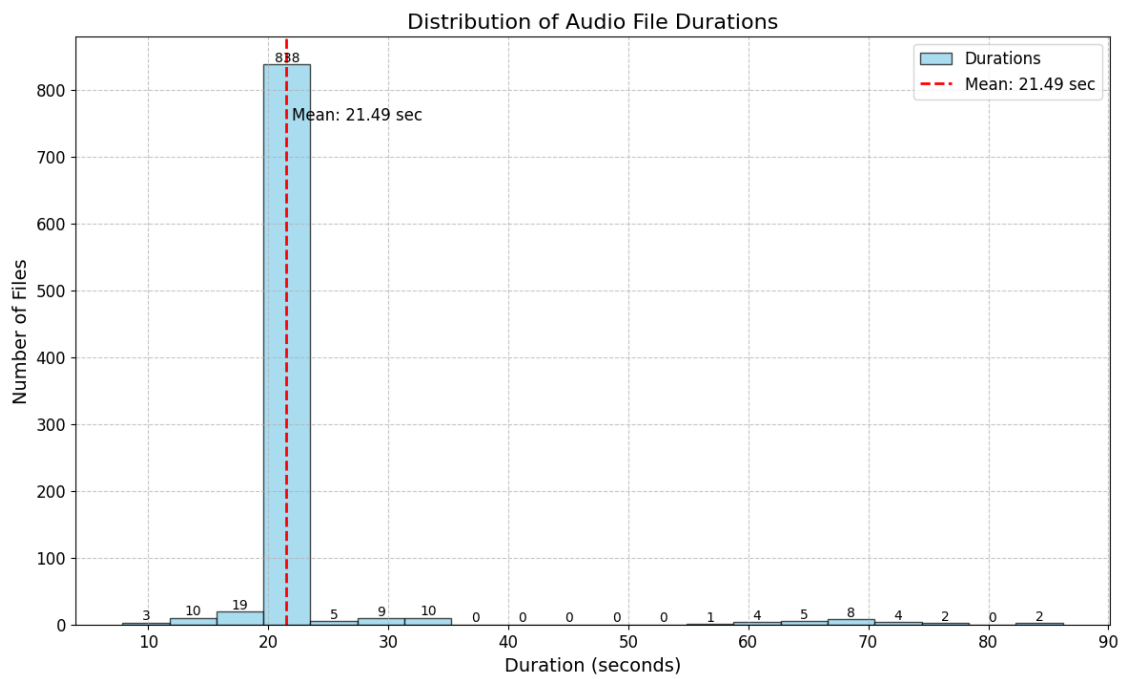
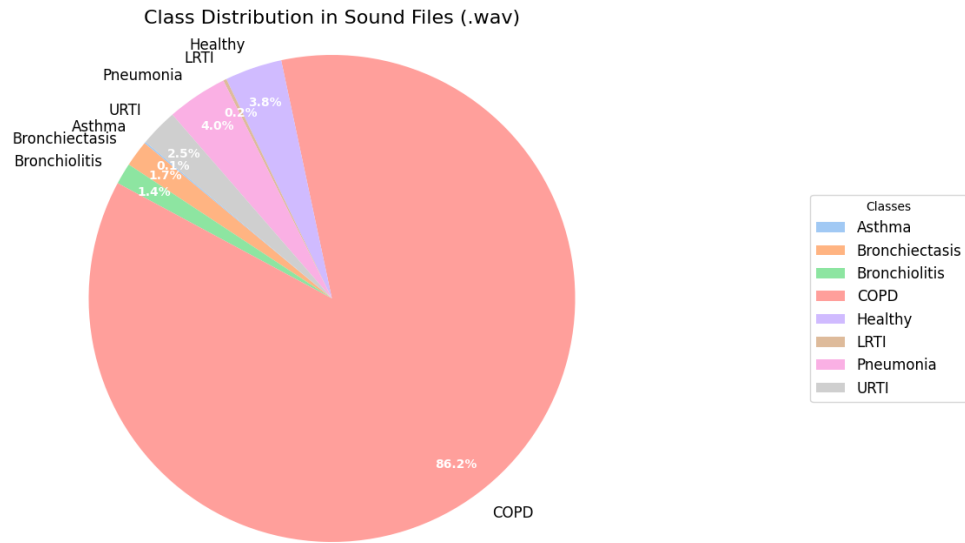
```

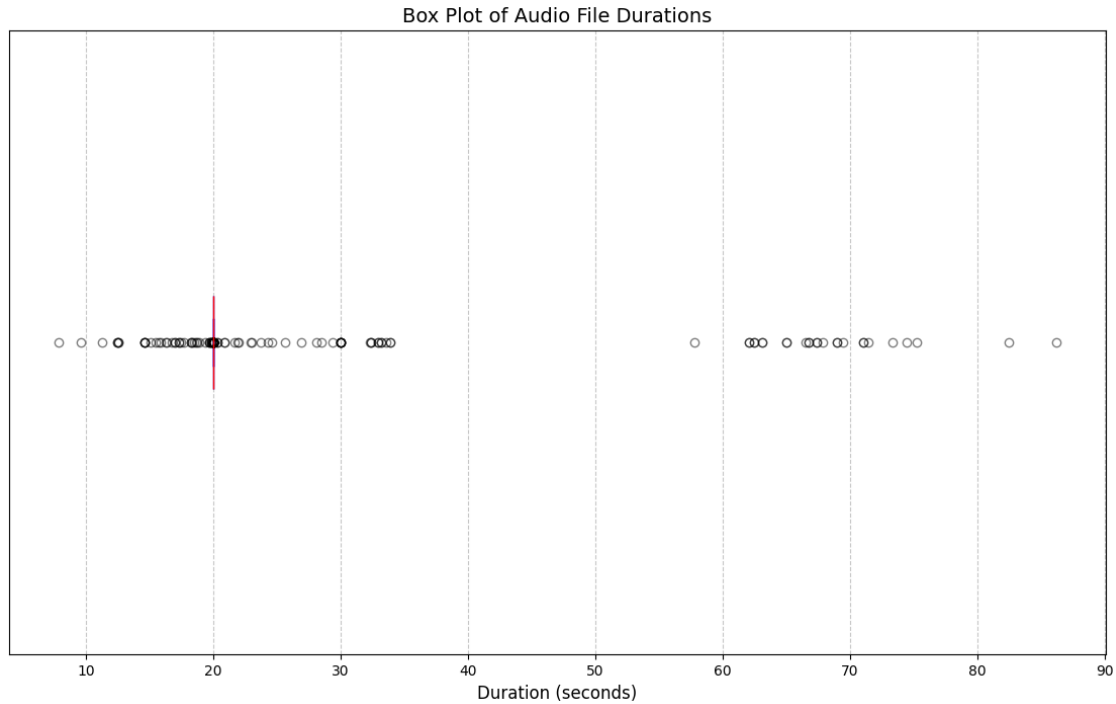
Classes	Count
Asthma	1
Bronchiectasis	16
Bronchiolitis	13
COPD	793
Healthy	35

LRTI	2
Pneumonia	37
URTI	23

Total Number of Classes: 8
Total Number of Patients: 920







Statistical Summary of Audio File Durations (in seconds)

	mean	std	min	max
0	21.49	8.31	7.86	86.2

1.5 Feature Extraction

`audio_features(filename)`: - Loads the audio file and retrieves the sample rate. - Computes the Short-Time Fourier Transform (STFT) of the audio. - Extracts various audio features, including:

- Mean Mel-frequency cepstral coefficients (MFCC).
- Mean chromagram from STFT.
- Mean mel-scaled spectrogram.
- Mean spectral contrast.
- Mean tonal centroid features (tonnetz).

- Concatenates all features into a single feature vector.

```
[7]: def audio_features(filename):
    try:
        # Load the audio file and get the sample rate
        sound, sample_rate = librosa.load(filename)
        # Compute the Short-Time Fourier Transform (STFT) of the sound
        stft = np.abs(librosa.stft(sound))

        # Compute the mean Mel-frequency cepstral coefficients (MFCC) over time
        mfccs = np.mean(librosa.feature.mfcc(y=sound, sr=sample_rate,
        ↪n_mfcc=40), axis=1)
        # Compute the mean chromagram from STFT
```

```

        chroma = np.mean(librosa.feature.chroma_stft(S=stft, sr=sample_rate),  

↪axis=1)  

        # Compute the mean mel-scaled spectrogram  

        mel = np.mean(librosa.feature.melspectrogram(y=sound, sr=sample_rate),  

↪axis=1)  

        # Compute the mean spectral contrast  

        contrast = np.mean(librosa.feature.spectral_contrast(S=stft,  

↪sr=sample_rate), axis=1)  

        # Compute the mean tonal centroid features (tonnetz)  

        tonnetz = np.mean(librosa.feature.tonnetz(y=librosa.effects.  

↪harmonic(sound), sr=sample_rate), axis=1)  

        # Concatenate all the features into a single feature vector  

        return np.concatenate((mfccs, chroma, mel, contrast, tonnetz))  

    except Exception as e:  

        raise RuntimeError(f"Error extracting audio features from file_  

↪{filename}: {e}")

```

data_points(PATH, diagnosis_file): - Initializes lists to store labels and features. - Maps diagnoses to numerical labels. - Loads diagnosis data. - Iterates over each diagnosis entry: - Appends the numerical label for the diagnosis to the labels list. - Extracts audio features and appends them to the features list. - Prints a completion message indicating the number of files processed. - Returns the labels and features as numpy arrays.

```

[8]: def data_points(PATH, diagnosis_file):  

    labels = [] # Initialize an empty list to store the labels  

    features = [] # Initialize an empty list to store the features  

    # Mapping of diagnosis to numerical labels  

    diagnosis_mapping = {  

        "COPD": 0, "Healthy": 1, "URTI": 2, "Bronchiectasis": 3,  

        "Pneumonia": 4, "Bronchiolitis": 5, "Asthma": 6, "LRTI": 7  

    }  

    try:  

        # Load diagnosis data from the given path and file  

        diagnoses = load_diagnosis_data(PATH, diagnosis_file)  

    except (ValueError, FileNotFoundError, RuntimeError) as e:  

        # Print error message if there is an issue with loading the diagnosis_  

↪data  

        print(f"Error loading diagnosis data: {e}")  

        return np.array([]), np.array([]) # Return empty arrays if an error_  

↪occurs  

    # Iterate over each diagnosis entry  

    for diagnosis in tqdm(diagnoses, desc="Extracting Features", unit="file"):  

        try:

```

```

        # Append the numerical label for the diagnosis to the labels list
        labels.append(diagnosis_mapping[diagnosis.diagnosis])
        # Extract and append the audio features to the features list
        features.append(audio_features(diagnosis.audio_file_path))
    except Exception as e:
        # Print a warning message if there is an issue with feature
↪extraction
        tqdm.write(f"\033[1;91mWARNING:\033[0m An issue was encountered_
↪while extracting features from file: {diagnosis.audio_file_path}. Error:_
↪{e}")

    # Print completion message indicating the number of files processed
    print(f'\n\033[1;92mFinished feature extraction from {len(features)}_
↪files\033[0m')
    # Return the labels and features as numpy arrays
    return np.array(labels), np.array(features)

```

1.6 Preprocessing

preprocessing(labels, images): - Removes entries with labels for Asthma (6) and LRTI (7). - Splits data into training and testing sets. - One-hot encodes the labels. - Reshapes the data for model input. - Returns the preprocessed training and testing sets.

Additional steps:

- Start timer: Measures the preprocessing time.
- Load data points: Calls the `data_points` function to retrieve labels and features.
- Preprocess data: Calls the `preprocessing` function to prepare the data.
- Calculate and print elapsed time: Displays the time taken for preprocessing in minutes and seconds.

Warning: Estimated Time of Completion (ETC): 22-25 minutes on a MacBook Air M2

```

[9]: def preprocessing(labels, images):
    # Remove Asthma (label 6) and LRTI (label 7)
    indices_to_remove = np.where((labels == 6) | (labels == 7))[0]
    images = np.delete(images, indices_to_remove, axis=0)
    labels = np.delete(labels, indices_to_remove, axis=0)

    # Split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(images, labels,
↪test_size=0.2, random_state=10)

    # One-hot encode the labels
    y_train = to_categorical(y_train)
    y_test = to_categorical(y_test)

    # Reshape data
    X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))

```



```

y_train = y_train.reshape((y_train.shape[0], 6))
X_test = X_test.reshape((X_test.shape[0], X_train.shape[1], 1))
y_test = y_test.reshape((y_test.shape[0], 6))

return X_train, X_test, y_train, y_test

# Start timer
start = timer()

# Load data points
labels, images = data_points(PATH, diagnosis_file)

# Preprocess data
X_train, X_test, y_train, y_test = preprocessing(labels, images)

end = timer()
elapsed_time = end - start

# Calculate minutes and remaining seconds
minutes = int(elapsed_time // 60)
seconds = int(elapsed_time % 60)

# Print time taken in a user-friendly format
print(f'Time taken: {minutes}:{seconds:02d} minutes ({elapsed_time:.0f}␣
↪seconds)')

```

Extracting Features: 0%| | 0/920 [00:00<?, ?file/s]

WARNING: An issue was encountered while extracting features from file:

../Respiratory_Sound_Database/audio_and_txt_files/101_1b1_A1_sc_Meditron.wav.

Error: 'Unknown'

WARNING: An issue was encountered while extracting features from file:

../Respiratory_Sound_Database/audio_and_txt_files/101_1b1_Pr_sc_Meditron.wav.

Error: 'Unknown'

Extracting Features: 0%| | 3/920 [00:10<53:39,

3.51s/file]/Users/afl/Library/Python/3.9/lib/python/site-

packages/librosa/core/pitch.py:101: UserWarning: Trying to estimate tuning from empty frequency set.

return pitch_tuning(

Extracting Features: 99%| | 908/920 [24:01<00:45, 3.83s/file]

/Users/afl/Library/Python/3.9/lib/python/site-

packages/librosa/core/pitch.py:101: UserWarning: Trying to estimate tuning from empty frequency set.

return pitch_tuning(

Extracting Features: 100%| | 920/920 [24:24<00:00, 1.59s/file]

Finished feature extraction from 918 files

Time taken: 24:24 minutes (1464 seconds)

1.6.1 Selecting a Random .wav File and Displaying its Mel-Spectrogram

`select_random_wav_file(PATH)`: - Lists all .wav files in the specified directory. - Randomly selects one .wav file. - Returns the full path of the selected file and its name.

`display_mel_spectrogram(file_path, file_name)`: - Loads the audio file at the given path. - Splices the first 5 seconds of the audio. - Computes the mel-spectrogram of the spliced audio. - Displays the mel-spectrogram with a color bar and appropriate labels.

Additional steps:

- Select a random .wav file: Calls the `select_random_wav_file` function to choose a file from the specified directory.
- Display mel-spectrogram: Calls the `display_mel_spectrogram` function to visualize the spectrogram of the selected audio file.
- Load necessary libraries: Imports `os`, `numpy`, `librosa`, and `matplotlib` libraries required for file handling, random selection, audio processing, and plotting.

```
[33]: # Function to randomly select a .wav file from a directory
def select_random_wav_file(PATH):
    files = [f for f in os.listdir(PATH) if f.endswith('.wav')]
    random_index = np.random.randint(0, len(files))
    selected_file = files[random_index]
    return os.path.join(PATH, selected_file), selected_file

# Function to display mel-spectrogram of the spliced audio file
def display_mel_spectrogram(file_path, file_name):
    # Load audio file
    y, sr = librosa.load(file_path, sr=None) # Set sr=None to load the native
    ↪ sampling rate

    # Splice audio (example: first 5 seconds)
    duration = 5 # seconds
    y_spliced = y[:sr * duration]

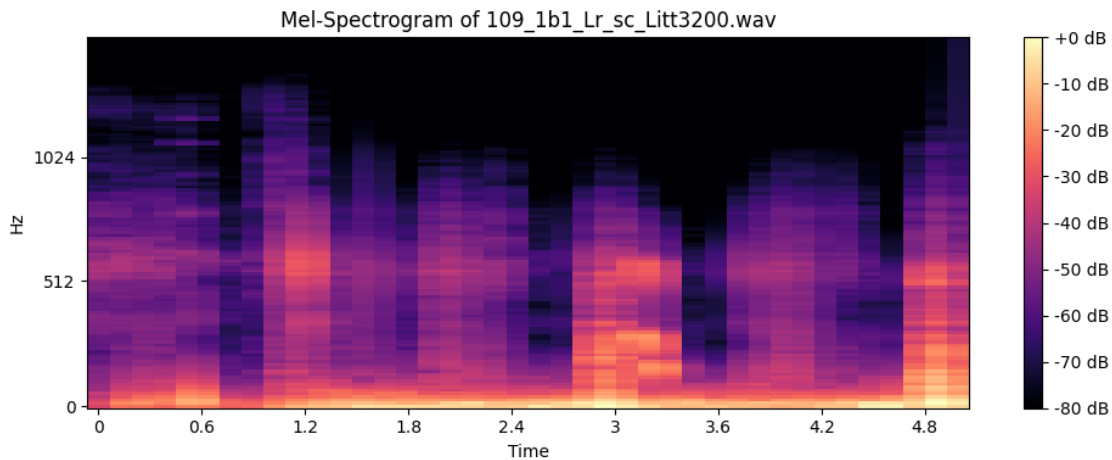
    # Compute mel-spectrogram
    mel_spectrogram = librosa.feature.melspectrogram(y=y_spliced, sr=sr,
    ↪ n_fft=2048, hop_length=512, n_mels=128)
    mel_spectrogram_db = librosa.power_to_db(mel_spectrogram, ref=np.max)

    # Display mel-spectrogram
    plt.figure(figsize=(10, 4))
    librosa.display.specshow(mel_spectrogram_db, sr=sr, hop_length=512,
    ↪ x_axis='time', y_axis='mel')
```

```
plt.colorbar(format='%+2.0f dB')
plt.title(f'Mel-Spectrogram of {file_name}')
plt.tight_layout()
plt.show()

# Select a random .wav file
random_wav_file, selected_file_name = select_random_wav_file(PATH)

# Display mel-spectrogram of the spliced audio file
display_mel_spectrogram(random_wav_file, selected_file_name)
```



1.7 Building the Models

1.7.1 Convolutional Neural Network (CNN)

The following functions are defined to build and train a CNN model:

- `build_model(input_shape, num_classes):`
 - Creates a CNN model with the following layers:
 - * Convolutional layers with ReLU activation.
 - * MaxPooling layer.
 - * Dropout layer.
 - * Flatten layer.
 - * Dense layers with ReLU and softmax activation.
 - Returns the constructed model.
- `compile_and_train_model(model, X_train, y_train, X_test, y_test, epochs=70, batch_size=200):`
 - Compiles the model with categorical cross-entropy loss and Adam optimizer.
 - Trains the model on the training data with validation on the test data.
 - Prints a completion message.
 - Returns the training history.

Additional steps:

- Define input shape and number of classes.
- Build the model: Calls `build_model` with the defined input shape and number of classes.
- Compile and train the model: Calls `compile_and_train_model` with the model and data.

```
[24]: def build_model(input_shape, num_classes):
    model = Sequential([
        Conv1D(64, kernel_size=5, activation='relu', input_shape=input_shape),
        Conv1D(128, kernel_size=5, activation='relu'),
        MaxPooling1D(pool_size=2),
        Conv1D(256, kernel_size=5, activation='relu'),
        Dropout(0.3),
        Flatten(),
        Dense(512, activation='relu'),
        Dense(num_classes, activation='softmax')
    ])
    return model

def compile_and_train_model(model, X_train, y_train, X_test, y_test, epochs=74,
    ↪batch_size=92):
    model.compile(loss='categorical_crossentropy', optimizer=Adam(),
    ↪metrics=['accuracy'])
    history = model.fit(X_train, y_train, validation_data=(X_test, y_test),
    ↪epochs=epochs, batch_size=batch_size, verbose=1)
    print("\033[92mCNN Model completed!\033[0m")
    return history

# Define input shape and number of classes
input_shape = (193, 1)
num_classes = 6

# Build the model
model = build_model(input_shape, num_classes)

# Compile and train the model
history = compile_and_train_model(model, X_train, y_train, X_test, y_test)
```

Epoch 1/74

8/8 2s 205ms/step -

accuracy: 0.5741 - loss: 6.5429 - val_accuracy: 0.7268 - val_loss: 0.9668

Epoch 2/74

8/8 2s 290ms/step -

accuracy: 0.7794 - loss: 1.0623 - val_accuracy: 0.8743 - val_loss: 0.5460

Epoch 3/74

8/8 2s 270ms/step -

accuracy: 0.8758 - loss: 0.8461 - val_accuracy: 0.8743 - val_loss: 0.5010

Epoch 4/74

8/8 2s 257ms/step -
accuracy: 0.8659 - loss: 0.5990 - val_accuracy: 0.8743 - val_loss: 0.4536
Epoch 5/74

8/8 2s 255ms/step -
accuracy: 0.8562 - loss: 0.5118 - val_accuracy: 0.8634 - val_loss: 0.4805
Epoch 6/74

8/8 2s 256ms/step -
accuracy: 0.8718 - loss: 0.4184 - val_accuracy: 0.8743 - val_loss: 0.4476
Epoch 7/74

8/8 2s 294ms/step -
accuracy: 0.8478 - loss: 0.4424 - val_accuracy: 0.8743 - val_loss: 0.4273
Epoch 8/74

8/8 3s 332ms/step -
accuracy: 0.8534 - loss: 0.4175 - val_accuracy: 0.8743 - val_loss: 0.3989
Epoch 9/74

8/8 2s 313ms/step -
accuracy: 0.8746 - loss: 0.3465 - val_accuracy: 0.8525 - val_loss: 0.3944
Epoch 10/74

8/8 2s 288ms/step -
accuracy: 0.8753 - loss: 0.3155 - val_accuracy: 0.8689 - val_loss: 0.3920
Epoch 11/74

8/8 2s 268ms/step -
accuracy: 0.8994 - loss: 0.2969 - val_accuracy: 0.8415 - val_loss: 0.4260
Epoch 12/74

8/8 2s 243ms/step -
accuracy: 0.9029 - loss: 0.3006 - val_accuracy: 0.8470 - val_loss: 0.3688
Epoch 13/74

8/8 2s 257ms/step -
accuracy: 0.8963 - loss: 0.2758 - val_accuracy: 0.8907 - val_loss: 0.3150
Epoch 14/74

8/8 2s 271ms/step -
accuracy: 0.9046 - loss: 0.2517 - val_accuracy: 0.8470 - val_loss: 0.3535
Epoch 15/74

8/8 2s 262ms/step -
accuracy: 0.9125 - loss: 0.2399 - val_accuracy: 0.8907 - val_loss: 0.3145
Epoch 16/74

8/8 2s 249ms/step -
accuracy: 0.9401 - loss: 0.1865 - val_accuracy: 0.8634 - val_loss: 0.3581
Epoch 17/74

8/8 2s 252ms/step -
accuracy: 0.9192 - loss: 0.1830 - val_accuracy: 0.8798 - val_loss: 0.3298
Epoch 18/74

8/8 2s 301ms/step -
accuracy: 0.9396 - loss: 0.1777 - val_accuracy: 0.8634 - val_loss: 0.3276
Epoch 19/74

8/8 3s 371ms/step -
accuracy: 0.9416 - loss: 0.1707 - val_accuracy: 0.8852 - val_loss: 0.3181
Epoch 20/74

8/8 3s 359ms/step -
accuracy: 0.9387 - loss: 0.1678 - val_accuracy: 0.8852 - val_loss: 0.3202
Epoch 21/74

8/8 3s 323ms/step -
accuracy: 0.9425 - loss: 0.1727 - val_accuracy: 0.8470 - val_loss: 0.3537
Epoch 22/74

8/8 4s 464ms/step -
accuracy: 0.9460 - loss: 0.1387 - val_accuracy: 0.8689 - val_loss: 0.3306
Epoch 23/74

8/8 3s 412ms/step -
accuracy: 0.9420 - loss: 0.1509 - val_accuracy: 0.8852 - val_loss: 0.3123
Epoch 24/74

8/8 4s 490ms/step -
accuracy: 0.9424 - loss: 0.1564 - val_accuracy: 0.8743 - val_loss: 0.3392
Epoch 25/74

8/8 3s 409ms/step -
accuracy: 0.9412 - loss: 0.1502 - val_accuracy: 0.8361 - val_loss: 0.3886
Epoch 26/74

8/8 3s 335ms/step -
accuracy: 0.9542 - loss: 0.1273 - val_accuracy: 0.8743 - val_loss: 0.3410
Epoch 27/74

8/8 2s 292ms/step -
accuracy: 0.9569 - loss: 0.1093 - val_accuracy: 0.8852 - val_loss: 0.3259
Epoch 28/74

8/8 2s 289ms/step -
accuracy: 0.9402 - loss: 0.1274 - val_accuracy: 0.8634 - val_loss: 0.3451
Epoch 29/74

8/8 2s 266ms/step -
accuracy: 0.9633 - loss: 0.0958 - val_accuracy: 0.8743 - val_loss: 0.3521
Epoch 30/74

8/8 2s 286ms/step -
accuracy: 0.9708 - loss: 0.0907 - val_accuracy: 0.8798 - val_loss: 0.3658
Epoch 31/74

8/8 2s 266ms/step -
accuracy: 0.9596 - loss: 0.0889 - val_accuracy: 0.8634 - val_loss: 0.3933
Epoch 32/74

8/8 2s 252ms/step -
accuracy: 0.9765 - loss: 0.0769 - val_accuracy: 0.8634 - val_loss: 0.3825
Epoch 33/74

8/8 3s 369ms/step -
accuracy: 0.9751 - loss: 0.0722 - val_accuracy: 0.8470 - val_loss: 0.3911
Epoch 34/74

8/8 3s 310ms/step -
accuracy: 0.9690 - loss: 0.0800 - val_accuracy: 0.8579 - val_loss: 0.3740
Epoch 35/74

8/8 2s 278ms/step -
accuracy: 0.9834 - loss: 0.0700 - val_accuracy: 0.8743 - val_loss: 0.3905
Epoch 36/74

8/8 3s 374ms/step -
 accuracy: 0.9654 - loss: 0.0878 - val_accuracy: 0.8852 - val_loss: 0.3575
 Epoch 37/74
 8/8 2s 266ms/step -
 accuracy: 0.9752 - loss: 0.0788 - val_accuracy: 0.8962 - val_loss: 0.3475
 Epoch 38/74
 8/8 2s 272ms/step -
 accuracy: 0.9834 - loss: 0.0499 - val_accuracy: 0.8689 - val_loss: 0.4048
 Epoch 39/74
 8/8 2s 267ms/step -
 accuracy: 0.9797 - loss: 0.0566 - val_accuracy: 0.8798 - val_loss: 0.3832
 Epoch 40/74
 8/8 3s 338ms/step -
 accuracy: 0.9829 - loss: 0.0520 - val_accuracy: 0.8743 - val_loss: 0.4401
 Epoch 41/74
 8/8 3s 349ms/step -
 accuracy: 0.9830 - loss: 0.0473 - val_accuracy: 0.8579 - val_loss: 0.4962
 Epoch 42/74
 8/8 2s 280ms/step -
 accuracy: 0.9777 - loss: 0.0539 - val_accuracy: 0.8743 - val_loss: 0.4903
 Epoch 43/74
 8/8 2s 302ms/step -
 accuracy: 0.9826 - loss: 0.0408 - val_accuracy: 0.8962 - val_loss: 0.4461
 Epoch 44/74
 8/8 2s 257ms/step -
 accuracy: 0.9918 - loss: 0.0390 - val_accuracy: 0.8907 - val_loss: 0.5081
 Epoch 45/74
 8/8 2s 260ms/step -
 accuracy: 0.9944 - loss: 0.0291 - val_accuracy: 0.8907 - val_loss: 0.4670
 Epoch 46/74
 8/8 2s 264ms/step -
 accuracy: 0.9856 - loss: 0.0333 - val_accuracy: 0.8852 - val_loss: 0.5020
 Epoch 47/74
 8/8 4s 498ms/step -
 accuracy: 0.9788 - loss: 0.0475 - val_accuracy: 0.8852 - val_loss: 0.4724
 Epoch 48/74
 8/8 4s 447ms/step -
 accuracy: 0.9907 - loss: 0.0330 - val_accuracy: 0.8689 - val_loss: 0.5114
 Epoch 49/74
 8/8 3s 307ms/step -
 accuracy: 0.9894 - loss: 0.0339 - val_accuracy: 0.8907 - val_loss: 0.5141
 Epoch 50/74
 8/8 2s 280ms/step -
 accuracy: 0.9940 - loss: 0.0248 - val_accuracy: 0.8962 - val_loss: 0.4631
 Epoch 51/74
 8/8 2s 273ms/step -
 accuracy: 0.9963 - loss: 0.0216 - val_accuracy: 0.8579 - val_loss: 0.5063
 Epoch 52/74

8/8 2s 277ms/step -
 accuracy: 0.9943 - loss: 0.0270 - val_accuracy: 0.8907 - val_loss: 0.5038
 Epoch 53/74
 8/8 3s 371ms/step -
 accuracy: 0.9838 - loss: 0.0318 - val_accuracy: 0.8907 - val_loss: 0.6233
 Epoch 54/74
 8/8 2s 308ms/step -
 accuracy: 0.9985 - loss: 0.0166 - val_accuracy: 0.8962 - val_loss: 0.4747
 Epoch 55/74
 8/8 2s 268ms/step -
 accuracy: 0.9993 - loss: 0.0165 - val_accuracy: 0.9016 - val_loss: 0.5038
 Epoch 56/74
 8/8 2s 313ms/step -
 accuracy: 0.9962 - loss: 0.0177 - val_accuracy: 0.8907 - val_loss: 0.5067
 Epoch 57/74
 8/8 2s 269ms/step -
 accuracy: 0.9990 - loss: 0.0148 - val_accuracy: 0.8907 - val_loss: 0.5111
 Epoch 58/74
 8/8 3s 322ms/step -
 accuracy: 0.9963 - loss: 0.0184 - val_accuracy: 0.8852 - val_loss: 0.6178
 Epoch 59/74
 8/8 3s 350ms/step -
 accuracy: 0.9988 - loss: 0.0160 - val_accuracy: 0.8962 - val_loss: 0.4783
 Epoch 60/74
 8/8 2s 265ms/step -
 accuracy: 0.9947 - loss: 0.0444 - val_accuracy: 0.8962 - val_loss: 0.4973
 Epoch 61/74
 8/8 2s 291ms/step -
 accuracy: 0.9945 - loss: 0.0521 - val_accuracy: 0.9016 - val_loss: 0.4074
 Epoch 62/74
 8/8 2s 289ms/step -
 accuracy: 0.9920 - loss: 0.0273 - val_accuracy: 0.8852 - val_loss: 0.5461
 Epoch 63/74
 8/8 3s 329ms/step -
 accuracy: 0.9885 - loss: 0.0331 - val_accuracy: 0.8743 - val_loss: 0.4684
 Epoch 64/74
 8/8 3s 305ms/step -
 accuracy: 0.9947 - loss: 0.0240 - val_accuracy: 0.8962 - val_loss: 0.4607
 Epoch 65/74
 8/8 2s 313ms/step -
 accuracy: 0.9915 - loss: 0.0241 - val_accuracy: 0.8689 - val_loss: 0.4884
 Epoch 66/74
 8/8 2s 316ms/step -
 accuracy: 0.9968 - loss: 0.0151 - val_accuracy: 0.8962 - val_loss: 0.4689
 Epoch 67/74
 8/8 3s 317ms/step -
 accuracy: 1.0000 - loss: 0.0071 - val_accuracy: 0.8907 - val_loss: 0.4932
 Epoch 68/74


```

8/8          3s 369ms/step -
accuracy: 0.9972 - loss: 0.0119 - val_accuracy: 0.8907 - val_loss: 0.5530
Epoch 69/74
8/8          3s 344ms/step -
accuracy: 0.9984 - loss: 0.0073 - val_accuracy: 0.9016 - val_loss: 0.5650
Epoch 70/74
8/8          3s 342ms/step -
accuracy: 1.0000 - loss: 0.0062 - val_accuracy: 0.9016 - val_loss: 0.5817
Epoch 71/74
8/8          3s 333ms/step -
accuracy: 1.0000 - loss: 0.0050 - val_accuracy: 0.8962 - val_loss: 0.5943
Epoch 72/74
8/8          3s 319ms/step -
accuracy: 1.0000 - loss: 0.0050 - val_accuracy: 0.8962 - val_loss: 0.5705
Epoch 73/74
8/8          2s 279ms/step -
accuracy: 1.0000 - loss: 0.0036 - val_accuracy: 0.8852 - val_loss: 0.5627
Epoch 74/74
8/8          2s 260ms/step -
accuracy: 1.0000 - loss: 0.0031 - val_accuracy: 0.8962 - val_loss: 0.5765
CNN Model completed!

```

1.7.2 Artificial Neural Network (ANN)

```

[12]: # Define the ANN model
def build_ann_model(input_shape, num_output_classes):
    ann_model = Sequential([
        Dense(256, activation='relu', input_shape=input_shape),
        Dropout(0.3),
        Dense(256, activation='relu'),
        Dropout(0.3),
        Dense(128, activation='relu'),
        Dropout(0.3),
        Dense(num_output_classes, activation='softmax')
    ])
    return ann_model

def compile_and_train_ann_model(ann_model, training_data, training_labels,
    ↪validation_data, validation_labels, num_epochs=70, batch_size_value=5):
    ann_model.compile(loss='categorical_crossentropy', optimizer=Adam(),
    ↪metrics=['accuracy'])
    training_history = ann_model.fit(training_data, training_labels,
    ↪validation_data=(validation_data, validation_labels), epochs=num_epochs,
    ↪batch_size=batch_size_value, verbose=1)
    print("\033[92mANN Model training completed!\033[0m")
    return training_history

```

```

# Define input shape and number of classes
input_shape_ann = (193,)
num_output_classes = 6

# Build the ANN model
ann_model = build_ann_model(input_shape_ann, num_output_classes)

# Compile and train the ANN model
training_history = compile_and_train_ann_model(ann_model, X_train, y_train,
↪X_test, y_test)

```

Epoch 1/70

```

/Users/afl/Library/Python/3.9/lib/python/site-
packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

147/147          1s 2ms/step -
accuracy: 0.7407 - loss: 13.6673 - val_accuracy: 0.8743 - val_loss: 1.7788

```

Epoch 2/70

```

147/147          0s 2ms/step -
accuracy: 0.7989 - loss: 2.7289 - val_accuracy: 0.8743 - val_loss: 0.5799

```

Epoch 3/70

```

147/147          0s 2ms/step -
accuracy: 0.7883 - loss: 1.7950 - val_accuracy: 0.8743 - val_loss: 0.5092

```

Epoch 4/70

```

147/147          0s 2ms/step -
accuracy: 0.7666 - loss: 1.0716 - val_accuracy: 0.8525 - val_loss: 0.7115

```

Epoch 5/70

```

147/147          0s 2ms/step -
accuracy: 0.8282 - loss: 0.6928 - val_accuracy: 0.8743 - val_loss: 0.6860

```

Epoch 6/70

```

147/147          0s 2ms/step -
accuracy: 0.7942 - loss: 0.8487 - val_accuracy: 0.8743 - val_loss: 0.5248

```

Epoch 7/70

```

147/147          0s 2ms/step -
accuracy: 0.8268 - loss: 0.6099 - val_accuracy: 0.8743 - val_loss: 0.4771

```

Epoch 8/70

```

147/147          0s 2ms/step -
accuracy: 0.8737 - loss: 0.5427 - val_accuracy: 0.8689 - val_loss: 0.4333

```

Epoch 9/70

```

147/147          0s 2ms/step -
accuracy: 0.8294 - loss: 0.6436 - val_accuracy: 0.8743 - val_loss: 0.4440

```

Epoch 10/70

```

147/147          0s 2ms/step -
accuracy: 0.8659 - loss: 0.4702 - val_accuracy: 0.8743 - val_loss: 0.4287

```

Epoch 11/70
147/147 0s 2ms/step -
accuracy: 0.8636 - loss: 0.5888 - val_accuracy: 0.8743 - val_loss: 0.3993

Epoch 12/70
147/147 0s 2ms/step -
accuracy: 0.8374 - loss: 0.5205 - val_accuracy: 0.8743 - val_loss: 0.5137

Epoch 13/70
147/147 0s 2ms/step -
accuracy: 0.8621 - loss: 0.4819 - val_accuracy: 0.8743 - val_loss: 0.4588

Epoch 14/70
147/147 0s 2ms/step -
accuracy: 0.8656 - loss: 0.4864 - val_accuracy: 0.8743 - val_loss: 0.4195

Epoch 15/70
147/147 0s 2ms/step -
accuracy: 0.8460 - loss: 0.4464 - val_accuracy: 0.8743 - val_loss: 0.4225

Epoch 16/70
147/147 0s 2ms/step -
accuracy: 0.8632 - loss: 0.4236 - val_accuracy: 0.8743 - val_loss: 0.4236

Epoch 17/70
147/147 0s 2ms/step -
accuracy: 0.8804 - loss: 0.3828 - val_accuracy: 0.8743 - val_loss: 0.4152

Epoch 18/70
147/147 0s 2ms/step -
accuracy: 0.8670 - loss: 0.4411 - val_accuracy: 0.8743 - val_loss: 0.4194

Epoch 19/70
147/147 0s 2ms/step -
accuracy: 0.8455 - loss: 0.4661 - val_accuracy: 0.8743 - val_loss: 0.4130

Epoch 20/70
147/147 0s 2ms/step -
accuracy: 0.8584 - loss: 0.4097 - val_accuracy: 0.8743 - val_loss: 0.4013

Epoch 21/70
147/147 0s 2ms/step -
accuracy: 0.8561 - loss: 0.3980 - val_accuracy: 0.8743 - val_loss: 0.4226

Epoch 22/70
147/147 0s 2ms/step -
accuracy: 0.8753 - loss: 0.3642 - val_accuracy: 0.8743 - val_loss: 0.4048

Epoch 23/70
147/147 0s 2ms/step -
accuracy: 0.8676 - loss: 0.3732 - val_accuracy: 0.8798 - val_loss: 0.3705

Epoch 24/70
147/147 0s 2ms/step -
accuracy: 0.8640 - loss: 0.3895 - val_accuracy: 0.8743 - val_loss: 0.3984

Epoch 25/70
147/147 0s 2ms/step -
accuracy: 0.8541 - loss: 0.4333 - val_accuracy: 0.8743 - val_loss: 0.3903

Epoch 26/70
147/147 0s 2ms/step -
accuracy: 0.8326 - loss: 0.4638 - val_accuracy: 0.8743 - val_loss: 0.3980

Epoch 27/70
147/147 0s 2ms/step -
accuracy: 0.8641 - loss: 0.3841 - val_accuracy: 0.8743 - val_loss: 0.4512

Epoch 28/70
147/147 0s 2ms/step -
accuracy: 0.8626 - loss: 0.3813 - val_accuracy: 0.8743 - val_loss: 0.3602

Epoch 29/70
147/147 0s 2ms/step -
accuracy: 0.8746 - loss: 0.4047 - val_accuracy: 0.8743 - val_loss: 0.3692

Epoch 30/70
147/147 0s 2ms/step -
accuracy: 0.8539 - loss: 0.4346 - val_accuracy: 0.8743 - val_loss: 0.3554

Epoch 31/70
147/147 0s 2ms/step -
accuracy: 0.8758 - loss: 0.4162 - val_accuracy: 0.8798 - val_loss: 0.3663

Epoch 32/70
147/147 0s 2ms/step -
accuracy: 0.8560 - loss: 0.3386 - val_accuracy: 0.8743 - val_loss: 0.4426

Epoch 33/70
147/147 0s 2ms/step -
accuracy: 0.9036 - loss: 0.3061 - val_accuracy: 0.8689 - val_loss: 0.4084

Epoch 34/70
147/147 0s 2ms/step -
accuracy: 0.8741 - loss: 0.3202 - val_accuracy: 0.8634 - val_loss: 0.4086

Epoch 35/70
147/147 0s 2ms/step -
accuracy: 0.8557 - loss: 0.3999 - val_accuracy: 0.8852 - val_loss: 0.3847

Epoch 36/70
147/147 0s 2ms/step -
accuracy: 0.8688 - loss: 0.3845 - val_accuracy: 0.8798 - val_loss: 0.3844

Epoch 37/70
147/147 0s 2ms/step -
accuracy: 0.8691 - loss: 0.3858 - val_accuracy: 0.8798 - val_loss: 0.3565

Epoch 38/70
147/147 0s 2ms/step -
accuracy: 0.9010 - loss: 0.2588 - val_accuracy: 0.8852 - val_loss: 0.3569

Epoch 39/70
147/147 0s 2ms/step -
accuracy: 0.8542 - loss: 0.3939 - val_accuracy: 0.8798 - val_loss: 0.4451

Epoch 40/70
147/147 0s 2ms/step -
accuracy: 0.8680 - loss: 0.3586 - val_accuracy: 0.8962 - val_loss: 0.3560

Epoch 41/70
147/147 0s 2ms/step -
accuracy: 0.8692 - loss: 0.3419 - val_accuracy: 0.8852 - val_loss: 0.4355

Epoch 42/70
147/147 0s 2ms/step -
accuracy: 0.8914 - loss: 0.3708 - val_accuracy: 0.8798 - val_loss: 0.4015

Epoch 43/70
147/147 0s 2ms/step -
accuracy: 0.8756 - loss: 0.3871 - val_accuracy: 0.8907 - val_loss: 0.3565
Epoch 44/70
147/147 0s 2ms/step -
accuracy: 0.8708 - loss: 0.3342 - val_accuracy: 0.8798 - val_loss: 0.3262
Epoch 45/70
147/147 0s 2ms/step -
accuracy: 0.8694 - loss: 0.3406 - val_accuracy: 0.8852 - val_loss: 0.3199
Epoch 46/70
147/147 0s 2ms/step -
accuracy: 0.9057 - loss: 0.2560 - val_accuracy: 0.8907 - val_loss: 0.3202
Epoch 47/70
147/147 0s 2ms/step -
accuracy: 0.8724 - loss: 0.3601 - val_accuracy: 0.8798 - val_loss: 0.3427
Epoch 48/70
147/147 0s 2ms/step -
accuracy: 0.8696 - loss: 0.3626 - val_accuracy: 0.8415 - val_loss: 0.3993
Epoch 49/70
147/147 0s 2ms/step -
accuracy: 0.8575 - loss: 0.3803 - val_accuracy: 0.8525 - val_loss: 0.3421
Epoch 50/70
147/147 0s 2ms/step -
accuracy: 0.8960 - loss: 0.3385 - val_accuracy: 0.8907 - val_loss: 0.3337
Epoch 51/70
147/147 0s 2ms/step -
accuracy: 0.8902 - loss: 0.2767 - val_accuracy: 0.9016 - val_loss: 0.3061
Epoch 52/70
147/147 0s 2ms/step -
accuracy: 0.8970 - loss: 0.2736 - val_accuracy: 0.9016 - val_loss: 0.3095
Epoch 53/70
147/147 0s 2ms/step -
accuracy: 0.8633 - loss: 0.4594 - val_accuracy: 0.8798 - val_loss: 0.3821
Epoch 54/70
147/147 0s 2ms/step -
accuracy: 0.8606 - loss: 0.3285 - val_accuracy: 0.8962 - val_loss: 0.4305
Epoch 55/70
147/147 0s 2ms/step -
accuracy: 0.8469 - loss: 0.3655 - val_accuracy: 0.8798 - val_loss: 0.3570
Epoch 56/70
147/147 0s 2ms/step -
accuracy: 0.8752 - loss: 0.3486 - val_accuracy: 0.8907 - val_loss: 0.3680
Epoch 57/70
147/147 0s 2ms/step -
accuracy: 0.8667 - loss: 0.3185 - val_accuracy: 0.8798 - val_loss: 0.3482
Epoch 58/70
147/147 0s 2ms/step -
accuracy: 0.8600 - loss: 0.3807 - val_accuracy: 0.8525 - val_loss: 0.3777

```

Epoch 59/70
147/147          0s 2ms/step -
accuracy: 0.8456 - loss: 0.3613 - val_accuracy: 0.8852 - val_loss: 0.3195
Epoch 60/70
147/147          0s 2ms/step -
accuracy: 0.8785 - loss: 0.3342 - val_accuracy: 0.8962 - val_loss: 0.3350
Epoch 61/70
147/147          0s 2ms/step -
accuracy: 0.8845 - loss: 0.3135 - val_accuracy: 0.8907 - val_loss: 0.3509
Epoch 62/70
147/147          0s 2ms/step -
accuracy: 0.8642 - loss: 0.3243 - val_accuracy: 0.8852 - val_loss: 0.3321
Epoch 63/70
147/147          0s 2ms/step -
accuracy: 0.9051 - loss: 0.2652 - val_accuracy: 0.9071 - val_loss: 0.3056
Epoch 64/70
147/147          0s 2ms/step -
accuracy: 0.8817 - loss: 0.2986 - val_accuracy: 0.9126 - val_loss: 0.3385
Epoch 65/70
147/147          0s 2ms/step -
accuracy: 0.8757 - loss: 0.3169 - val_accuracy: 0.8962 - val_loss: 0.4042
Epoch 66/70
147/147          0s 2ms/step -
accuracy: 0.8962 - loss: 0.2879 - val_accuracy: 0.8415 - val_loss: 0.4125
Epoch 67/70
147/147          0s 2ms/step -
accuracy: 0.8901 - loss: 0.3035 - val_accuracy: 0.9016 - val_loss: 0.3322
Epoch 68/70
147/147          0s 2ms/step -
accuracy: 0.8645 - loss: 0.3064 - val_accuracy: 0.9071 - val_loss: 0.3754
Epoch 69/70
147/147          0s 2ms/step -
accuracy: 0.8802 - loss: 0.2733 - val_accuracy: 0.9016 - val_loss: 0.2740
Epoch 70/70
147/147          0s 2ms/step -
accuracy: 0.8765 - loss: 0.3395 - val_accuracy: 0.9235 - val_loss: 0.2931
ANN Model training completed!

```

1.8 Evaluation - CNN

1.8.1 Plotting Accuracy & Loss

The model is evaluated on the test set, and accuracy and loss are printed in green. Learning curves for accuracy and loss are plotted to visualize model performance over epochs. The training and validation accuracy and loss values are displayed with markers, grid lines, titles, and labels for clarity.

```
[25]: # Evaluate the model
score = model.evaluate(X_test, y_test, batch_size=60, verbose=0)

# Print accuracy and loss in green
accuracy = score[1]
loss = score[0]

print(f'\033[1;92mAccuracy: {accuracy:.0%}\033[0m')

time.sleep(1)

print(f"\033[1;92mLoss: {loss:.4f}\033[0m\n")

time.sleep(1)

# Plot learning curves for accuracy and loss
fig, axes = plt.subplots(1, 2, figsize=(15, 5))


# Plot training & validation accuracy values
axes[0].plot(history.history['accuracy'], label='Training Accuracy', marker='o')
axes[0].plot(history.history['val_accuracy'], label='Validation Accuracy',
             ↪marker='x')
axes[0].set_title('Model Accuracy Over Epochs', fontsize=16)
axes[0].set_ylabel('Accuracy', fontsize=14)
axes[0].set_xlabel('Epoch', fontsize=14)
axes[0].legend(loc='upper left', fontsize=12)
axes[0].grid(True)

# Plot training & validation loss values
axes[1].plot(history.history['loss'], label='Training Loss', marker='o')
axes[1].plot(history.history['val_loss'], label='Validation Loss', marker='x')
axes[1].set_title('Model Loss Over Epochs', fontsize=16)
axes[1].set_ylabel('Loss', fontsize=14)
axes[1].set_xlabel('Epoch', fontsize=14)
axes[1].legend(loc='upper left', fontsize=12)
axes[1].grid(True)

plt.tight_layout()
plt.show()
```

Accuracy: 90%

Loss: 0.5765



1.8.2 Plotting Confusion Matrix

The confusion matrix and classification report are generated to evaluate model predictions against true labels. The confusion matrix is then displayed with percentage annotations using a heatmap.

```
[26]: # Define matrix index
matrix_index = ["COPD", "Healthy", "URTI", "Bronchiectasis", "Pneumoina",
               ↪ "Bronchiolitis"]

# Make predictions
preds = model.predict(X_test)
class_preds = np.argmax(preds, axis=1) # Predicted classes
y_test_class = np.argmax(y_test, axis=1) # True classes

# Compute confusion matrix
cm = confusion_matrix(y_test_class, class_preds)
print(classification_report(y_test_class, class_preds,
    ↪ target_names=matrix_index))

# Calculate percentage values for the confusion matrix
cm_sum = np.sum(cm, axis=1, keepdims=True)
cm_perc = cm / cm_sum.astype(float) * 100

# Create annotations for the confusion matrix
annot = np.empty_like(cm).astype(str)
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        c = cm[i, j]
        p = cm_perc[i, j]
        if i == j:
            annot[i, j] = f'{p:.1f}%\n{c}/{cm_sum[i][0]}'
        elif c == 0:
            annot[i, j] = ''
        else:
```



```

annot[i, j] = f'{p:.1f}%\n{c}'

# Display the confusion matrix
df_cm = pd.DataFrame(cm, index=matrix_index, columns=matrix_index)
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'

plt.figure(figsize=(12, 8))
heatmap = sn.heatmap(df_cm, annot=annot, fmt='', cmap='YlGnBu', cbar=False,
                    linewidths=.5, linecolor='black')
heatmap.set_title('Confusion Matrix', fontsize=18)
heatmap.set_xlabel('Predicted Label', fontsize=14)
heatmap.set_ylabel('True Label', fontsize=14)
plt.yticks(rotation=0, fontsize=12)
plt.xticks(rotation=45, fontsize=12)
plt.tight_layout()
plt.show()

```

WARNING:tensorflow:5 out of the last 13 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x31ad7ec10> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

```

6/6          0s 25ms/step
      precision    recall  f1-score   support

      COPD          0.97      0.97      0.97        160
      Healthy        0.50      0.14      0.22          7
      URTI           0.00      0.00      0.00          2
Bronchiectasis      0.60      0.60      0.60          5
      Pneumoina      0.60      0.75      0.67          4
Bronchiolitis       0.33      0.20      0.25          5

      accuracy                   0.90        183
      macro avg          0.50      0.44      0.45        183
      weighted avg       0.90      0.90      0.90        183

```

		Confusion Matrix					
True Label	COPD	97.5% 156/160		0.6% 1	0.6% 1	1.2% 2	
	Healthy	14.3% 1	14.3% 1/7	42.9% 3			28.6% 2
	URTI	50.0% 1	50.0% 1	0.0% 0/2			
	Bronchiectasis	40.0% 2			60.0% 3/5		
	Pneumonia	25.0% 1				75.0% 3/4	
	Bronchiolitis			60.0% 3	20.0% 1		20.0% 1/5
		COPD	Healthy	URTI	Bronchiectasis	Pneumonia	Bronchiolitis
		Predicted Label					

1.8.3 Plotting ROC & AUC

ROC curves are computed and plotted for each class to evaluate the model's performance. The curves and their respective AUC values are displayed in a plot with distinct colors for each class.

```
[27]: # Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
n_classes = y_test.shape[1]

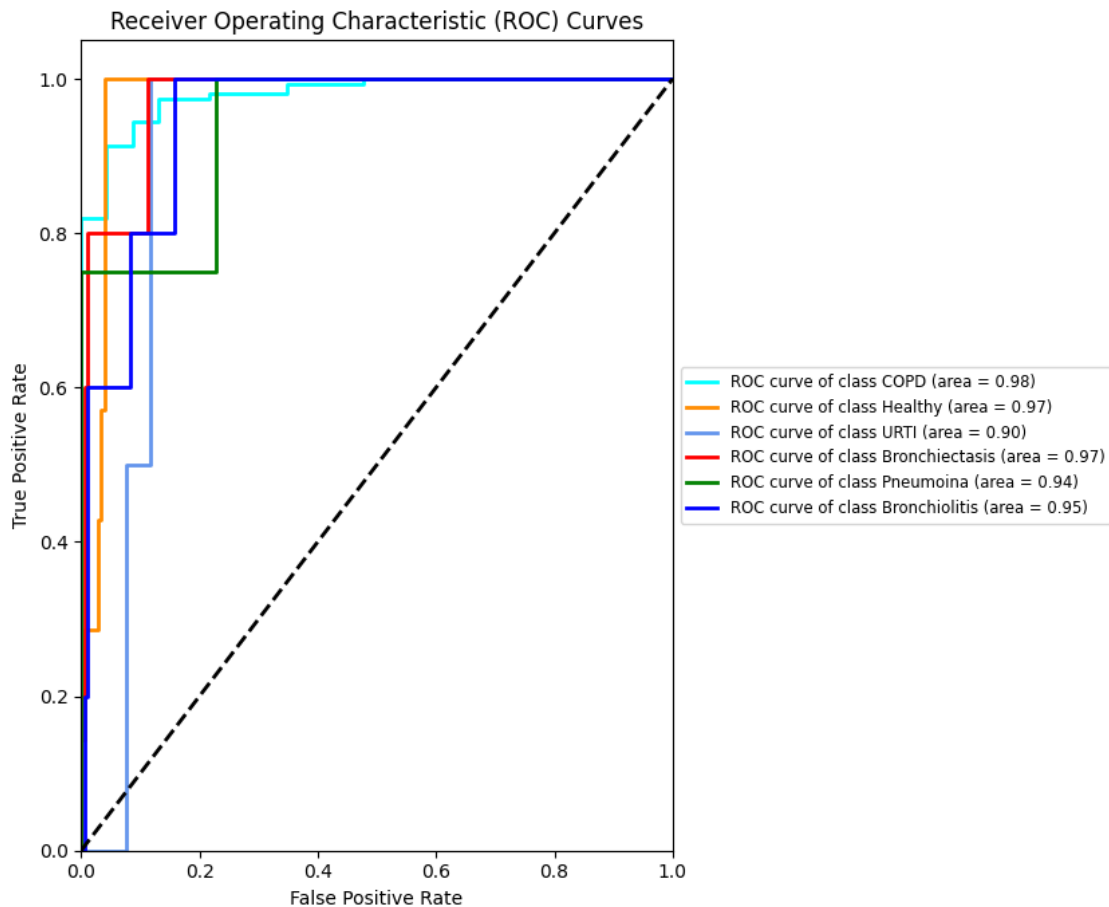
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], preds[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC curves
plt.figure(figsize=(10, 7))
colors = cycle(['aqua', 'darkorange', 'cornflowerblue', 'red', 'green', 'blue'])
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label=f'ROC curve of class {matrix_index[i]} (area = {roc_auc[i]:.2f})')
```

```

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curves')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), fontsize='small') #
    ↳Adjust legend position
plt.tight_layout(rect=[0, 0, 0.85, 1]) # Make room on the right for the legend
plt.show()

```



1.8.4 Plotting Precision-Recall Curves

Precision-recall curves are computed and plotted for each class to further evaluate model performance. The curves are displayed with distinct colors and labels for each class, including a legend and grid for clarity.

```
[28]: n_classes = 6
colors = ['cyan', 'orange', 'blue', 'red', 'green', 'purple']
matrix_index = ['COPD', 'Healthy', 'URTI', 'Bronchiectasis', 'Pneumonia', '
↳Bronchiolitis']

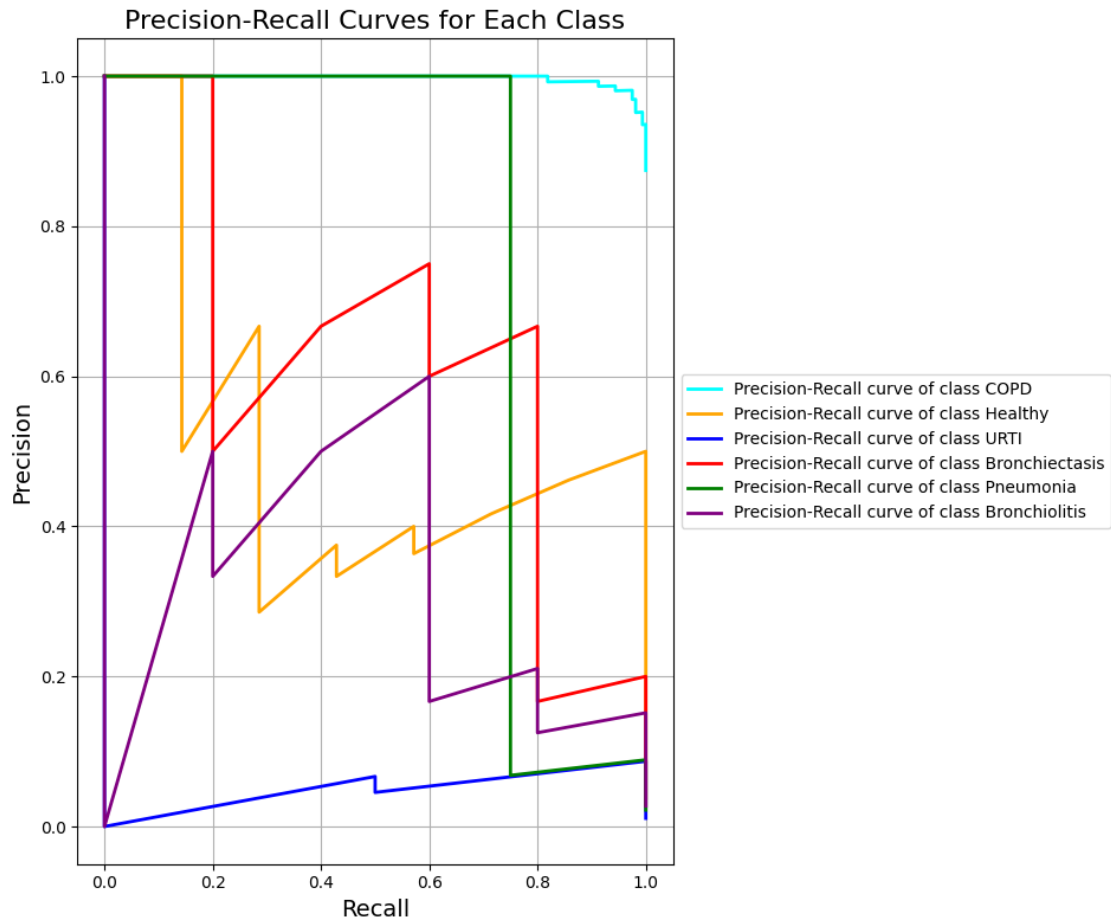
# Compute precision-recall curve for each class
precision = dict()
recall = dict()
for i in range(n_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test[:, i], preds[:, i])
↳i])

# Plot Precision-Recall curves
plt.figure(figsize=(12, 8))
for i, color in zip(range(n_classes), colors):
    plt.plot(recall[i], precision[i], color=color, lw=2,
              label=f'Precision-Recall curve of class {matrix_index[i]}')

# Adding grid, titles, and labels
plt.grid(True)
plt.xlabel('Recall', fontsize=14)
plt.ylabel('Precision', fontsize=14)
plt.title('Precision-Recall Curves for Each Class', fontsize=16)
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), fontsize='medium') #
↳Adjust legend position

# Adjust layout to make room for the legend
plt.tight_layout(rect=[0, 0, 0.8, 1])

plt.show()
```



1.9 Evaluation - ANN

1.9.1 Plotting Accuracy & Loss

```
[17]: # Evaluate the ANN model
score = ann_model.evaluate(X_test, y_test, batch_size=200, verbose=0)

# Print accuracy and loss in green
accuracy2 = score[1]
loss2 = score[0]

print(f'\033[1;92mAccuracy: {accuracy2:.0%}\033[0m')

time.sleep(1)

print(f"\033[1;92mLoss: {loss2:.4f}\033[0m\n")

time.sleep(1)
```

```

# Plot learning curves for accuracy and loss
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# Plot training & validation accuracy values
axes[0].plot(training_history.history['accuracy'], label='Training Accuracy',
             ↪marker='o')
axes[0].plot(training_history.history['val_accuracy'], label='Validation
             ↪Accuracy', marker='x')
axes[0].set_title('Model Accuracy Over Epochs', fontsize=16)
axes[0].set_ylabel('Accuracy', fontsize=14)
axes[0].set_xlabel('Epoch', fontsize=14)
axes[0].legend(loc='upper left', fontsize=12)
axes[0].grid(True)


# Plot training & validation loss values
axes[1].plot(training_history.history['loss'], label='Training Loss',
             ↪marker='o')
axes[1].plot(training_history.history['val_loss'], label='Validation Loss',
             ↪marker='x')
axes[1].set_title('Model Loss Over Epochs', fontsize=16)
axes[1].set_ylabel('Loss', fontsize=14)
axes[1].set_xlabel('Epoch', fontsize=14)
axes[1].legend(loc='upper left', fontsize=12)
axes[1].grid(True)

plt.tight_layout()
plt.show()

```

Accuracy: 92%

Loss: 0.2931



1.9.2 Plotting Confusion Matrix

```
[18]: # Define matrix index
matrix_index = ["COPD", "Healthy", "URTI", "Bronchiectasis", "Pneumonia",
               ↪ "Bronchiolitis"]

# Make predictions
preds = ann_model.predict(X_test)
class_preds = np.argmax(preds, axis=1) # Predicted classes
y_test_class = np.argmax(y_test, axis=1) # True classes

# Compute confusion matrix
cm = confusion_matrix(y_test_class, class_preds)
print(classification_report(y_test_class, class_preds,
    ↪ target_names=matrix_index))

# Calculate percentage values for the confusion matrix
cm_sum = np.sum(cm, axis=1, keepdims=True)
cm_perc = cm / cm_sum.astype(float) * 100

# Create annotations for the confusion matrix
annot = np.empty_like(cm).astype(str)
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        c = cm[i, j]
        p = cm_perc[i, j]
        if i == j:
            annot[i, j] = f'{p:.1f}%\n{c}/{cm_sum[i][0]}'
        elif c == 0:
            annot[i, j] = ''
        else:
            annot[i, j] = f'{p:.1f}%\n{c}'

# Display the confusion matrix
df_cm = pd.DataFrame(cm, index=matrix_index, columns=matrix_index)
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'

plt.figure(figsize=(12, 8))
heatmap = sn.heatmap(df_cm, annot=annot, fmt='', cmap='YlGnBu', cbar=False,
    ↪ linewidths=.5, linecolor='black')
heatmap.set_title('Confusion Matrix', fontsize=18)
heatmap.set_xlabel('Predicted Label', fontsize=14)
heatmap.set_ylabel('True Label', fontsize=14)
plt.yticks(rotation=0, fontsize=12)
plt.xticks(rotation=45, fontsize=12)
plt.tight_layout()
```

```
plt.show()
```

```
6/6          0s 5ms/step
          precision    recall  f1-score   support

      COPD          0.96      0.99      0.98        160
    Healthy          0.50      1.00      0.67         7
      URTI          0.00      0.00      0.00         2
Bronchiectasis        1.00      0.40      0.57         5
    Pneumonia        1.00      0.25      0.40         4
Bronchiolitis        0.00      0.00      0.00         5

 accuracy                   0.92        183
  macro avg          0.58      0.44      0.44        183
  weighted avg          0.91      0.92      0.90        183
```

```
/Users/afl/Library/Python/3.9/lib/python/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/Users/afl/Library/Python/3.9/lib/python/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
/Users/afl/Library/Python/3.9/lib/python/site-
packages/sklearn/metrics/_classification.py:1509: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```


		Confusion Matrix					
True Label	COPD	99.4% 159/160	0.6% 1				
	Healthy		100.0% 7/7				
	URTI		100.0% 2	0.0% 0/2			
	Bronchiectasis	60.0% 3			40.0% 2/5		
	Pneumonia	75.0% 3				25.0% 1/4	
	Bronchiolitis	20.0% 1	80.0% 4				0.0% 0/5
		COPD	Healthy	URTI	Bronchiectasis	Pneumonia	Bronchiolitis
		Predicted Label					

1.9.3 Plotting ROC & AUC

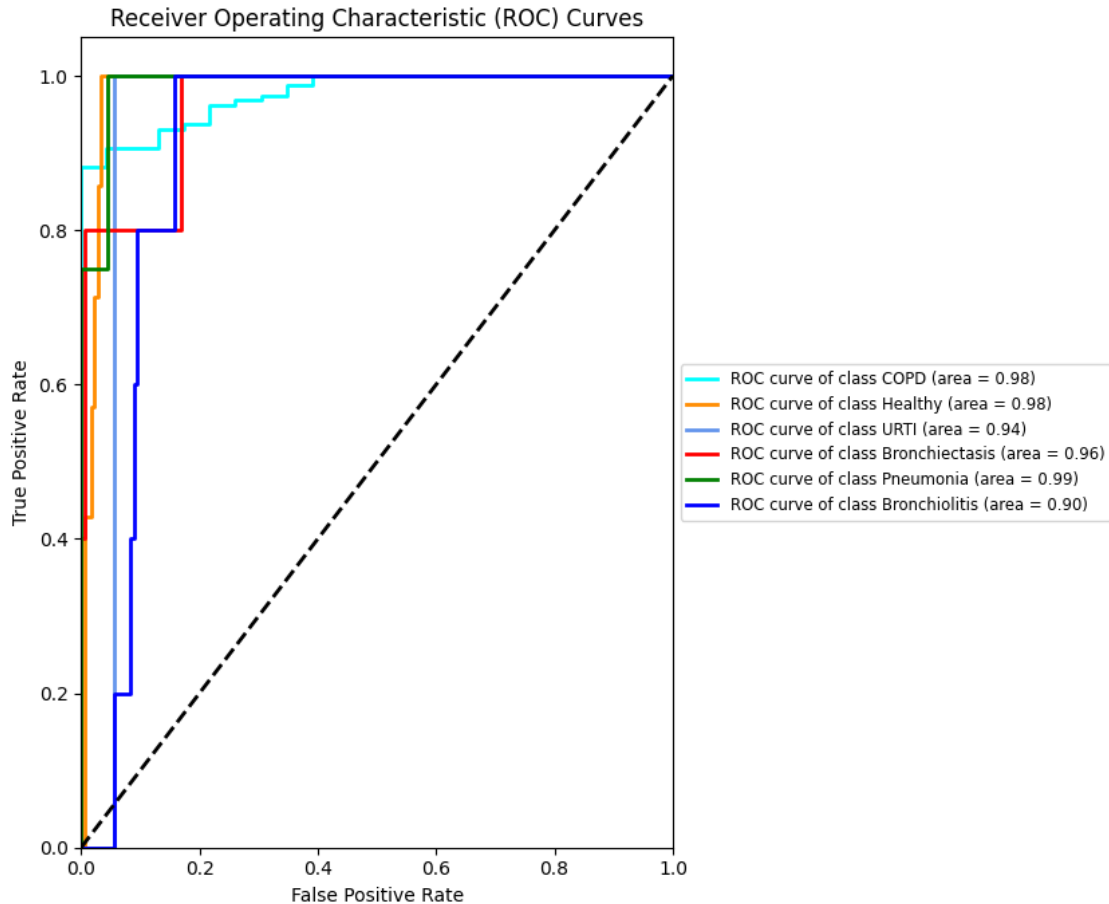
```
[19]: # Compute ROC curve and ROC area for each class
false_positive_rate = dict()
true_positive_rate = dict()
roc_auc_score = dict()
num_classes = y_test.shape[1]

for i in range(num_classes):
    false_positive_rate[i], true_positive_rate[i], _ = roc_curve(y_test[:, i],
↳ preds[:, i])
    roc_auc_score[i] = auc(false_positive_rate[i], true_positive_rate[i])

# Plot ROC curves
plt.figure(figsize=(10, 7))
colors = cycle(['aqua', 'darkorange', 'cornflowerblue', 'red', 'green', 'blue'])
for i, color in zip(range(num_classes), colors):
    plt.plot(false_positive_rate[i], true_positive_rate[i], color=color, lw=2,
        label=f'ROC curve of class {matrix_index[i]} (area =
↳ {roc_auc_score[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
```

```
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curves')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), fontsize='small') #
    ↪ Adjust legend position
plt.tight_layout(rect=[0, 0, 0.85, 1]) # Make room on the right for the legend
plt.show()
```



1.9.4 Plotting Precision-Recall Curves

```
[20]: # Define the number of classes, class names, and colors
num_classes = 6
class_names = ['COPD', 'Healthy', 'URTI', 'Bronchiectasis', 'Pneumonia',
    ↪ 'Bronchiolitis']
colors = ['cyan', 'orange', 'blue', 'red', 'green', 'purple']

# Compute precision-recall curve for each class
```

```

precision_dict = {}
recall_dict = {}
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test[:, i], preds[:, i])

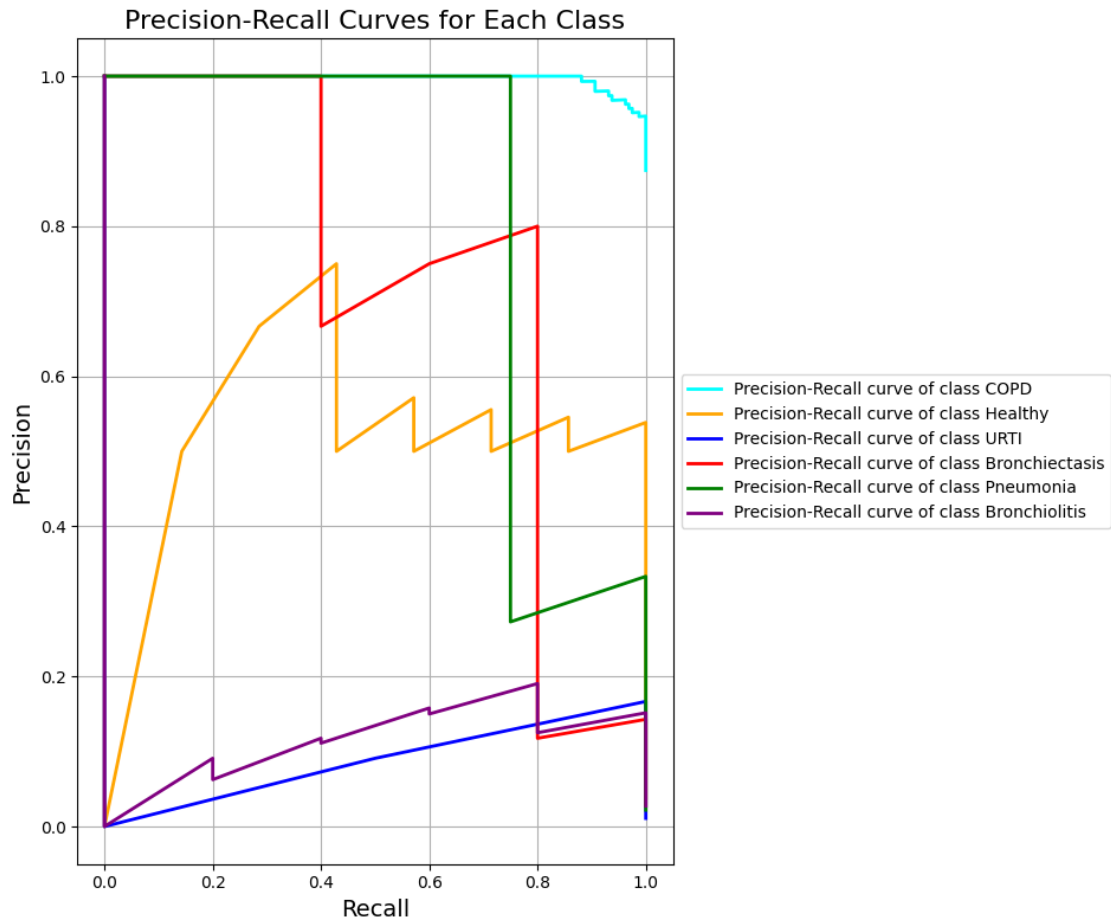
# Plot Precision-Recall curves
plt.figure(figsize=(12, 8))
for i, color in zip(range(num_classes), colors):
    plt.plot(recall[i], precision[i], color=color, lw=2,
             label=f'Precision-Recall curve of class {class_names[i]}')

# Adding grid, titles, and labels
plt.grid(True)
plt.xlabel('Recall', fontsize=14)
plt.ylabel('Precision', fontsize=14)
plt.title('Precision-Recall Curves for Each Class', fontsize=16)
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5), fontsize='medium') # Adjust legend position

# Adjust layout to make room for the legend
plt.tight_layout(rect=[0, 0, 0.8, 1])

plt.show()

```



1.10 Comparison of the two Models

```
[29]: # Data for Model 1 (ANN Model)
ann_accuracy = accuracy2 # Using accuracy obtained from ANN model evaluation
ann_loss = loss2 # Using loss obtained from ANN model evaluation

# Data for Model 2 (CNN Model)
cnn_accuracy = accuracy # Using accuracy obtained from model evaluation
cnn_loss = loss # Using loss obtained from model evaluation

# Define slightly darker pastel colors
pastel_blue = '#7EC8E3'
pastel_green = '#7ECFA2'

# Plotting
labels = ['ANN Model', 'CNN Model']
accuracies = [ann_accuracy, cnn_accuracy]
losses = [ann_loss, cnn_loss]
```

```

x = range(len(labels))
width = 0.35

fig, ax = plt.subplots(figsize=(14, 10))

# Plot accuracy bars
ax.bar(x, accuracies, width, label='Accuracy', color=pastel_blue, alpha=0.7)


# Plot loss bars
ax.bar([i + width for i in x], losses, width, label='Loss', color=pastel_green,
      ↪alpha=0.7)

# Attach a text label above each bar displaying its height
for i, (accuracy, loss) in enumerate(zip(accuracies, losses)):
    ax.annotate(f'{accuracy:.3f}', xy=(i, accuracy), xytext=(0, 3),
      ↪textcoords="offset points", ha='center', va='bottom')
    ax.annotate(f'{loss:.3f}', xy=(i + width, loss), xytext=(0, 3),
      ↪textcoords="offset points", ha='center', va='bottom')

ax.set_ylabel('Scores')
ax.set_title('Comparison of Model Evaluation Metrics')
ax.set_xticks([i + width / 2 for i in x])
ax.set_xticklabels(labels)
ax.legend()
ax.grid(True, linestyle='--', alpha=0.6)

plt.show()

```



1.11 Summary and Conclusion

This notebook performed an in-depth analysis and modelling of respiratory sound data, which included the following key steps: 1. **Data Loading and Preprocessing:** - Audio files and diagnosis data were loaded and preprocessed. - Relevant audio features were extracted for model training.

2. Data Visualization:

- Various plots were generated to visualize class distributions, audio file durations, and model performance.

3. Model Building and Evaluation:

- A Convolutional Neural Network (CNN) was built and trained to classify respiratory diseases.
- An Artificial Neural Network (ANN) was built and trained to classify respiratory diseases.
- The models were evaluated using accuracy, loss, confusion matrix, ROC curves, and precision-recall curves.

4. Results:

- The models showed promising results in classifying different respiratory conditions.
- Visualizations provided insights into the model's performance and helped identify areas for improvement.